

---

# ContextOS

## A Plain-Language Guide

How your context operating system works — explained without code, jargon, or technical structures.

---

Prepared June 7, 2026

For: the ContextOS owner

# Contents

---

|                                          |    |
|------------------------------------------|----|
| 1. What ContextOS Is                     | 4  |
| 2. The Three Big Ideas                   | 5  |
| 3. Meet the Cast                         | 6  |
| 4. A Tour of Your Workspace              | 7  |
| 5. How Work Actually Flows               | 9  |
| 6. Staying in Control — Safety by Design | 10 |
| 7. Seeing What Happened — Observability  | 11 |
| 8. Everyday Things You Can Do            | 12 |
| 9. Glossary                              | 13 |
| 10. Getting the Most Out of ContextOS    | 15 |



# 1. What ContextOS Is

---

ContextOS is a private control room where you put AI assistants to work — safely. It keeps the knowledge your assistants need, gives them the ability to take real actions on your behalf, and keeps you firmly in charge of everything they do.

Most AI tools are a single chat box: you ask, it answers, and the moment ends there. ContextOS is different. It remembers what matters to you, it can actually do things in the real world (look things up, contact services, send email), and it keeps a complete, reviewable record of every action. Think of it less like a chatbot and more like a small, trustworthy operations team that works for you around the clock.

The name says it plainly: it is an operating system for context. Just as your computer's operating system quietly manages files, programs, and permissions so your apps can run, ContextOS manages knowledge, assistants, tools, and permissions so your AI helpers can get useful work done without going off the rails.

## In one sentence

ContextOS gives your AI assistants memory, the ability to act, and guardrails — so they can do real work for you while you stay in control.

## 2. The Three Big Ideas

---

Everything in ContextOS rests on three simple ideas: Context, Capabilities, and Control.

### **Context — what your assistants know**

Context is the knowledge your assistants draw on: notes you've saved, facts about your work, preferences, and things they've learned from past tasks. Instead of re-explaining yourself every time, you teach ContextOS once and your assistants carry that understanding into every conversation and every task.

### **Capabilities — what your assistants can do**

Capabilities are the real actions your assistants can take. Rather than only talking, they can search and read the web, connect to outside services, and send messages — but only the specific actions you've allowed. Each capability is like a labelled tool in a toolbox: clearly named, with a known purpose.

### **Control — how you stay in charge**

Control is the set of guardrails that keeps everything safe: approvals for anything sensitive, lists of who is allowed to interact with your assistants, risk labels on every tool, and a complete history of what happened. Nothing important happens behind your back.

#### **Why all three matter together**

Knowledge without the ability to act is just a notebook. The ability to act without control is risky. ContextOS combines all three so your assistants are helpful, capable, and safe at the same time.

## 3. Meet the Cast

---

A handful of characters do all the work in ContextOS. Once you know who's who, the whole system makes sense.

### You — the owner

This is your private workspace. You set everything up, decide what your assistants are allowed to do, and approve anything sensitive. Everyone and everything else exists to serve your goals.

### Agents — your specialised assistants

Agents are AI workers you configure for particular jobs — for example, a research assistant, a writing assistant, or an inbox helper. Each one has its own role, its own knowledge, and its own set of permitted tools. You can have as many as you like.

### The Bot — your always-on personal assistant

The Bot is a single, general assistant you can reach from outside the app — by Telegram message or by email. It's the friendly front door: you ask it for something in plain language, and it either answers directly or quietly hands the work to the right agent.

### Models — the 'brains' you plug in

A model is the underlying AI engine that powers an assistant's thinking. ContextOS lets you choose which model each assistant uses, so you can balance quality, speed, and cost without changing anything else.

### Allowed contacts — who may interact

You decide which people and which systems are permitted to talk to your assistants. An email from a stranger, for instance, won't be acted on unless you've allowed that sender.

## 4. A Tour of Your Workspace

---

ContextOS is organised into clearly labelled areas in a sidebar. Here is what each one is for, in everyday terms.

### Dashboard

Your at-a-glance home screen. It shows what's happening right now — recent activity, anything waiting for you, and quick notices such as when an assistant has built a new tool for you.

### Agents

Where you create and manage your specialised assistants. You give each one a name and a role, choose its brain, and decide which tools and knowledge it can use.

### Model Endpoints

Where you connect the AI brains. You can plug in different models and pick which assistant uses which one — handy for matching the right level of capability and cost to each job.

### MCP Servers — your toolbox

This is where your assistants' tools live. A 'server' is simply a bundle of related tools that connect to an outside service. Remarkably, your assistant can build new tools for you here on request — wiring up a web service so future tasks can use it. Each tool can be tested on demand and earns a 'Verified' badge once it's confirmed working, so you can trust it.

### Accounts

Where connections to outside services and their sign-in details are stored safely, so your assistants can use those services without you re-entering credentials each time.

### Chat

A familiar chat screen for talking directly to your assistants inside the app. You can pick which assistant answers, ask questions, and kick off tasks — and you'll see live updates as work progresses.

### Telegram

Connects your assistant to Telegram so you can message it from your phone like any other contact, and get replies and results wherever you are.

### Email

Gives your assistant its own email address. You can email it tasks and it can email you (and allowed contacts) back — including properly formatted messages and attachments. You control exactly who is allowed to email it.

### Intents

Saved goals or playbooks — recurring things you want done a certain way. Instead of re-describing a task each time, you capture it once as an intent and reuse it.

## Runs

A 'run' is a single piece of work an assistant carries out, start to finish. This area lists every run so you can watch progress, open one up, and see exactly what the assistant did, step by step.

## Approvals

Your inbox for decisions. When an assistant wants to do something sensitive, it pauses and asks here. Nothing risky proceeds until you say yes.

## Artifacts

The things your assistants produce and the materials they work with — documents, results, and saved outputs — kept in one place.

## Memory

The long-term knowledge your assistants keep: facts, preferences, and lessons from past work. You can review and curate what they remember.

## Bot

The settings for your always-on personal assistant — how it behaves, what it's allowed to do, and how it handles messages that arrive by Telegram and email.

## Traces

A detailed, behind-the-scenes view of how a task unfolded — every step, decision, and tool use — for when you want the full story of what happened.

## Evaluations

Where the quality of your assistants' work is checked and scored over time, so you can tell whether things are getting better and spot what needs improvement.

## Audit Log

A tamper-evident diary of everything that happened: who did what, when, and why. This is your accountability record.

## Remote Access

Lets approved outside tools and assistants connect to your ContextOS securely, so your setup can plug into the wider world on your terms.

## API Keys

Secure passes that let trusted programs connect to your workspace. You can create them, see them, and switch them off at any time.

## Settings

The general controls for your whole workspace — the overall preferences that tie everything together.



## 5. How Work Actually Flows

---

The best way to understand ContextOS is to follow a single request from start to finish. Here's a realistic example.

Suppose you message your assistant: “Check the latest pricing on these three competitor websites and email me a short summary.” Here's what happens behind the scenes:

1. Your assistant understands the request. The Bot reads your message (from Telegram, email, or the in-app chat) and works out what you actually want.
2. It starts a run. The work becomes a tracked task — a 'run' — so you can watch it unfold and review it later.
3. It gathers the right context. The assistant pulls in what it already knows about you and the job, so it doesn't have to ask obvious questions.
4. It uses tools to act. It safely visits the three websites and reads the pricing — using only the web tools you've allowed, each one already tested and verified.
5. It pauses for anything sensitive. If a step needs your sign-off, it stops and asks you in Approvals. Nothing risky proceeds without your yes.
6. It produces and delivers the result. It writes the summary and emails it to you, formatted neatly.
7. Everything is recorded. Each step is saved to the run, the trace, and the audit log, so you can see exactly what was done and how.

### Notice the safety checkpoints

In that single request there were several guardrails: only allowed tools were used, only verified tools were trusted, anything sensitive needed your approval, and every step was logged. That combination is what makes it safe to let assistants act on your behalf.

## 6. Staying in Control — Safety by Design

---

ContextOS assumes that assistants should be powerful but never unaccountable. Safety isn't an add-on; it's built into how everything works.

### Risk levels on every tool

Each tool carries a risk label. Gentle, read-only actions can run freely, while anything with real-world consequences is treated with extra caution and may require your approval.

### Human approvals

For sensitive actions, the assistant stops and asks you first. You review the request in plain language and decide. This is your brake pedal.

### Allowed contacts only

You control exactly who can interact with your assistants. A message from an unknown sender isn't acted on unless you've allowed them — protecting you from strangers trying to give your assistant instructions.

### Safe web access

When assistants browse or search the web, they do so through a protected channel designed to prevent them from reaching places they shouldn't, keeping your private network and systems out of reach.

### Verified tools that stay honest

A tool only shows a 'Verified' badge once it's been tested and confirmed working. Importantly, if a tool's setup later changes — for example, its sign-in details are updated — that badge is automatically cleared until the tool is tested again. This means a 'Verified' label never lies about a tool that may have just been altered.

### A complete audit log

Every meaningful action is written to a diary you can review: what happened, who or what triggered it, and when. If you ever need to understand or prove what occurred, the record is there.

### Secrets kept safe

Passwords, keys, and other sensitive details are stored securely and never casually exposed — not in conversations, not in logs, not on screen.

## 7. Seeing What Happened — Observability

---

ContextOS is built so you're never in the dark. You can always see what your assistants are doing and how well they're doing it.

### Runs and traces

Every task is recorded as a run, and each run can be opened to reveal a step-by-step trace — the full story of what the assistant did. If something goes wrong, you can see exactly where and why, instead of guessing.

### Live work versus practice runs

Some actions happen for real ('live'), while others can be simulated safely as a dry run ('practice') — useful for checking that a tool works without causing real-world effects. ContextOS keeps track of which is which, so you always know whether something genuinely happened or was just rehearsed.

### Evaluations

Over time, the quality of your assistants' work can be scored and reviewed. This helps you spot what's working well, catch what isn't, and steadily improve how your assistants perform.

#### The point of all this

Trust comes from transparency. Because you can always inspect what happened and measure how well it went, you can confidently hand more work to your assistants.

## 8. Everyday Things You Can Do

---

Here are practical, real-world ways people put ContextOS to work.

- Message your assistant on Telegram to look something up and report back, wherever you are.
- Email a task to your assistant and get a tidy, formatted reply — with attachments if needed.
- Have a research assistant gather information from the web and summarise it for you.
- Set up recurring playbooks (intents) for tasks you do often, so they run the same way every time.
- Let an assistant draft and send routine messages, while you approve anything that matters.
- Connect an outside service once, then let your assistants use it safely for future tasks.
- Ask your assistant to build a new web tool for a service you use, then test and verify it.
- Review a complete record of everything your assistants have done, any time you like.

### Start small

You don't have to set everything up at once. Many people begin with a single assistant and one channel — say, Telegram — and grow from there as they get comfortable.

## 9. Glossary

---

Plain-language definitions of the terms you'll come across.

### **Agent**

A specialised AI assistant you configure for a particular kind of work, with its own role, knowledge, and permitted tools.

### **Bot**

Your single, always-on personal assistant, reachable from outside the app by Telegram and email.

### **Context**

The knowledge your assistants draw on — saved notes, facts, preferences, and lessons from past tasks.

### **Memory**

The long-term knowledge your assistants keep and reuse across many tasks.

### **Intent**

A saved goal or playbook for a task you want done a certain way, so you can reuse it instead of re-explaining.

### **Run**

A single piece of work an assistant carries out from start to finish, recorded so you can review it.

### **Trace**

The detailed, step-by-step record of how a particular run unfolded.

### **Approval**

A pause where the assistant asks for your permission before doing something sensitive.

### **Tool**

A specific action an assistant can take, such as searching the web or contacting a service.

### **MCP Server**

A bundle of related tools that connects to an outside service. Your assistant can build new ones on request.

### **Verified**

A badge showing a tool has been tested and confirmed working. It clears automatically if the tool's setup later changes.

### **Model / Model Endpoint**

The underlying AI 'brain' that powers an assistant's thinking, and the connection you set up to use it.

### **Account (Linked Account)**

A stored, secure connection to an outside service so assistants can use it without re-entering details.

### **Allowed contact**

A person or system you've permitted to interact with your assistants.

### **Risk level**

A label on each tool indicating how cautious to be — gentle actions run freely, consequential ones need care.

**Audit Log**

A complete, reviewable diary of everything that happened in your workspace.

**Evaluation**

A check and score of how well your assistants' work meets your standards over time.

**Artifact**

Something your assistants produce or work with — a document, result, or saved output.

**API Key**

A secure pass that lets a trusted program connect to your workspace; you can switch it off any time.

**Remote Access**

A secure way for approved outside tools and assistants to connect to your ContextOS.

## 10. Getting the Most Out of ContextOS

---

A few habits will help you get reliable, trustworthy results.

- Teach it once. Save the facts and preferences you'd otherwise repeat, so every assistant benefits.
- Be deliberate about tools. Only allow the actions you actually want, and keep risky ones behind approvals.
- Test before you trust. Use the on-demand test to verify a tool, and re-test after any change.
- Mind your allowed contacts. Keep the list of who can reach your assistants tight and intentional.
- Use the record. When in doubt about what happened, open the run, trace, or audit log — the answer is there.
- Grow gradually. Add assistants, channels, and tools as your confidence and needs increase.

### The big takeaway

ContextOS lets you delegate real work to AI assistants the way you'd delegate to a capable team: give them what they need to know, let them act, but keep clear rules, approvals, and a full record. That balance of capability and control is the whole point.